

adult-census-income-prediction

June 7, 2023

1 Adult census income prediction

1.1 Problem statement :

- Problem Statement: The Goal is to predict whether a person has an income of more than 50K a year or not. This is basically a binary classification problem where a person is classified into the >50K group or <=50K group.

1.1.1 import required libraries of python

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn import model_selection, preprocessing, tree
from sklearn.metrics import \
    confusion_matrix, classification_report, accuracy_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
import seaborn as sns
from sklearn.tree import export_graphviz
from sklearn.tree import DecisionTreeClassifier
```

1.1.2 import dataset :

kaggle link :<https://www.kaggle.com/datasets/overload10/adult-census-dataset>

```
[2]: df = pd.read_csv('adult.csv')
```

```
[3]: df
```

```
[3]:
```

	age	workclass	fnlwgt	education	education-num	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	
...	

32556	27	Private	257302	Assoc-acdm	12
32557	40	Private	154374	HS-grad	9
32558	58	Private	151910	HS-grad	9
32559	22	Private	201490	HS-grad	9
32560	52	Self-emp-inc	287927	HS-grad	9

	marital-status	occupation	relationship	race	\
0	Never-married	Adm-clerical	Not-in-family	White	
1	Married-civ-spouse	Exec-managerial	Husband	White	
2	Divorced	Handlers-cleaners	Not-in-family	White	
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	
4	Married-civ-spouse	Prof-specialty	Wife	Black	
...	
32556	Married-civ-spouse	Tech-support	Wife	White	
32557	Married-civ-spouse	Machine-op-inspct	Husband	White	
32558	Widowed	Adm-clerical	Unmarried	White	
32559	Never-married	Adm-clerical	Own-child	White	
32560	Married-civ-spouse	Exec-managerial	Wife	White	

	sex	capital-gain	capital-loss	hours-per-week	country	\
0	Male	2174	0	40	United-States	
1	Male	0	0	13	United-States	
2	Male	0	0	40	United-States	
3	Male	0	0	40	United-States	
4	Female	0	0	40	Cuba	
...	
32556	Female	0	0	38	United-States	
32557	Male	0	0	40	United-States	
32558	Female	0	0	40	United-States	
32559	Male	0	0	20	United-States	
32560	Female	15024	0	40	United-States	

	salary
0	<=50K
1	<=50K
2	<=50K
3	<=50K
4	<=50K
...	...
32556	<=50K
32557	>50K
32558	<=50K
32559	<=50K
32560	>50K

[32561 rows x 15 columns]

1.2 EDA

1.2.1 The info() method prints information about the DataFrame.

The information contains the number of columns, column labels, column data types, memory usage, range index, and the number of cells in each column (non-null values).

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32561 non-null  int64
1   workclass             32560 non-null  object
2   fnlwgt               32561 non-null  int64
3   education             32561 non-null  object
4   education-num         32561 non-null  int64
5   marital-status       32561 non-null  object
6   occupation            32561 non-null  object
7   relationship         32561 non-null  object
8   race                 32561 non-null  object
9   sex                  32561 non-null  object
10  capital-gain          32561 non-null  int64
11  capital-loss          32561 non-null  int64
12  hours-per-week        32561 non-null  int64
13  country               32561 non-null  object
14  salary                32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

```
[5]: df.describe()
```

```
[5]:
```

	age	fnlwgt	education-num	capital-gain	capital-loss	\
count	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	
mean	38.581647	1.897784e+05	10.080679	1077.648844	87.303830	
std	13.640433	1.055500e+05	2.572720	7385.292085	402.960219	
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000	
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	

	hours-per-week
count	32561.000000
mean	40.437456
std	12.347429
min	1.000000

25%	40.000000
50%	40.000000
75%	45.000000
max	99.000000

[6]: df

```
[6]:
```

	age	workclass	fnlwgt	education	education-num	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	
...	
32556	27	Private	257302	Assoc-acdm	12	
32557	40	Private	154374	HS-grad	9	
32558	58	Private	151910	HS-grad	9	
32559	22	Private	201490	HS-grad	9	
32560	52	Self-emp-inc	287927	HS-grad	9	

	marital-status	occupation	relationship	race	\
0	Never-married	Adm-clerical	Not-in-family	White	
1	Married-civ-spouse	Exec-managerial	Husband	White	
2	Divorced	Handlers-cleaners	Not-in-family	White	
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	
4	Married-civ-spouse	Prof-specialty	Wife	Black	
...	
32556	Married-civ-spouse	Tech-support	Wife	White	
32557	Married-civ-spouse	Machine-op-inspct	Husband	White	
32558	Widowed	Adm-clerical	Unmarried	White	
32559	Never-married	Adm-clerical	Own-child	White	
32560	Married-civ-spouse	Exec-managerial	Wife	White	

	sex	capital-gain	capital-loss	hours-per-week	country	\
0	Male	2174	0	40	United-States	
1	Male	0	0	13	United-States	
2	Male	0	0	40	United-States	
3	Male	0	0	40	United-States	
4	Female	0	0	40	Cuba	
...	
32556	Female	0	0	38	United-States	
32557	Male	0	0	40	United-States	
32558	Female	0	0	40	United-States	
32559	Male	0	0	20	United-States	
32560	Female	15024	0	40	United-States	

salary

```

0      <=50K
1      <=50K
2      <=50K
3      <=50K
4      <=50K
...
32556  <=50K
32557  >50K
32558  <=50K
32559  <=50K
32560  >50K

```

[32561 rows x 15 columns]

```
[7]: print(df['salary'].value_counts())
```

```

<=50K    24720
>50K      7841
Name: salary, dtype: int64

```

```
[8]: print(df['sex'].value_counts())
```

```

Male      21790
Female    10771
Name: sex, dtype: int64

```

```
[9]: a = [df['country'].value_counts()]
a
```

```
[9]: [ United-States    29170
      Mexico          643
      ?              583
      Philippines     198
      Germany         137
      Canada          121
      Puerto-Rico     114
      El-Salvador     106
      India           100
      Cuba            95
      England         90
      Jamaica         81
      South           80
      China           75
      Italy           73
      Dominican-Republic 70
      Vietnam         67
      Guatemala       64

```

Japan	62
Poland	60
Columbia	59
Taiwan	51
Haiti	44
Iran	43
Portugal	37
Nicaragua	34
Peru	31
France	29
Greece	29
Ecuador	28
Ireland	24
Hong	20
Cambodia	19
Trinidad&Tobago	19
Laos	18
Thailand	18
Yugoslavia	16
Outlying-US(Guam-USVI-etc)	14
Honduras	13
Hungary	13
Scotland	12
Holand-Netherlands	1

Name: country, dtype: int64]

```
[10]: df['workclass'].value_counts()
```


```
[10]: Private          22696
Self-emp-not-inc     2541
Local-gov            2093
?                    1835
State-gov            1298
Self-emp-inc         1116
Federal-gov          960
Without-pay          14
Never-worked          7
Name: workclass, dtype: int64
```

1.3 Data visualization


1.4 Age

```
[11]: plt.figure(figsize=(8,4))
sns.histplot(df['age'],color='green',bins=15)
plt.tight_layout()
plt.grid(True)
```

```
plt.title('Age distribution')  
plt.show()
```



```
[12]: sns.histplot(x=df['age'],hue=df['salary'],color='green',bins=15)  
plt.tight_layout()  
plt.title('Age distribution')  
plt.show()
```



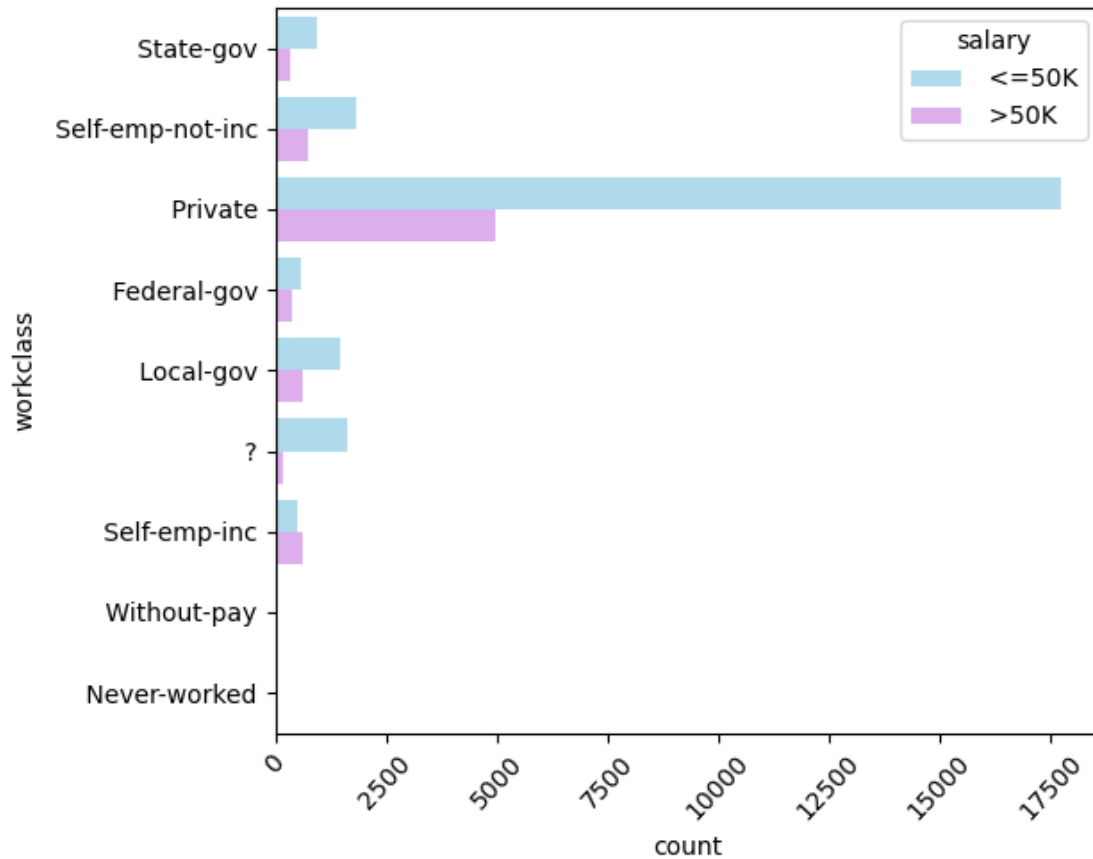
From the graph we can see that in the age group 0-20 there isn't any entry of salary greater than 50k, same goes with the group greater than 75 years.

1.4.1 work class

```
[13]: df['workclass'].unique()
```

```
[13]: array([' State-gov', ' Self-emp-not-inc', ' Private', ' Federal-gov',
        ' Local-gov', ' ?', ' Self-emp-inc', nan, ' Without-pay',
        ' Never-worked'], dtype=object)
```

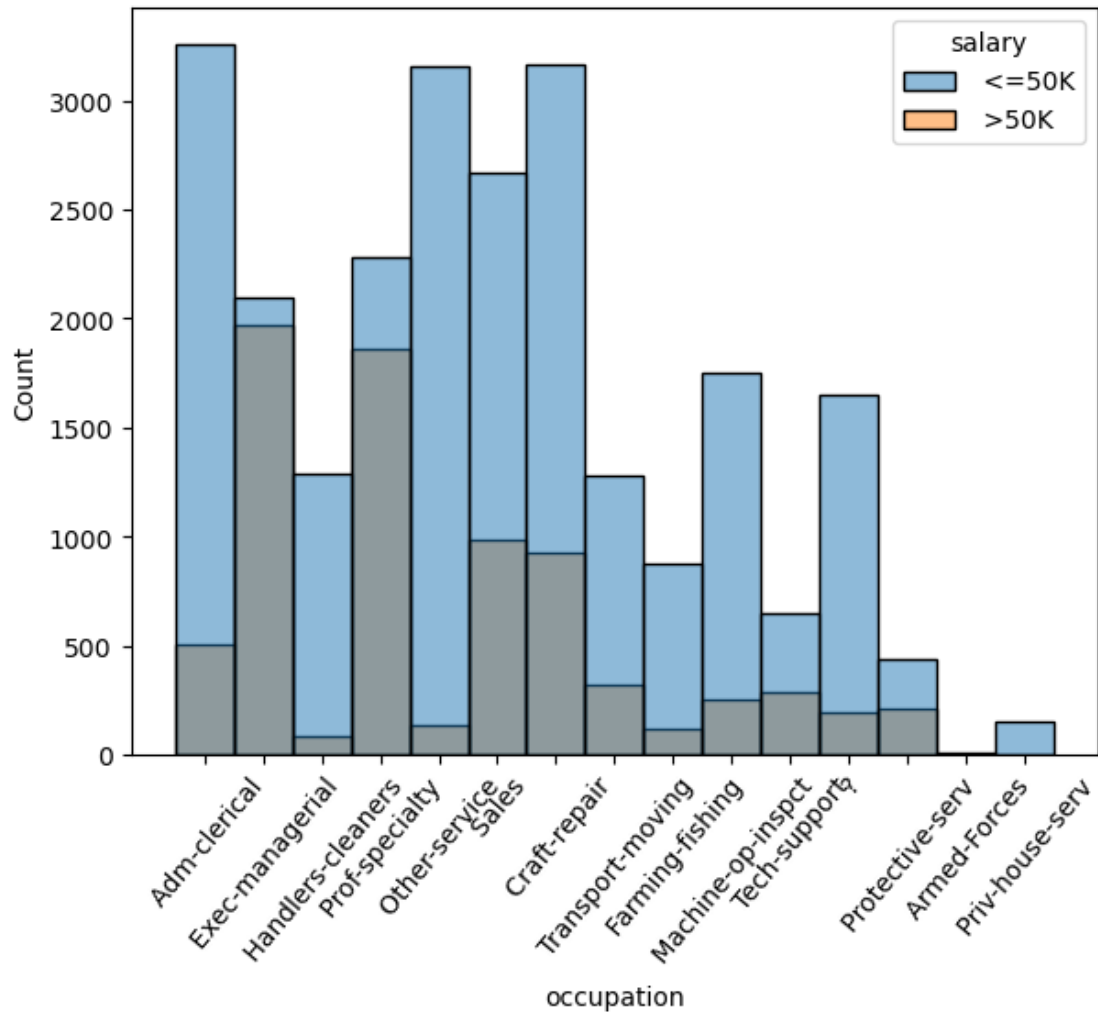
```
[17]: sns.countplot(y=df['workclass'], hue = df['salary'],
        palette=['#a4def5', '#e1a4f5'])
plt.tight_layout()
plt.xticks(rotation = 47)
plt.show()
```

- The majority of the individuals work in the private sector. The probabilities of making above 50,000 are similar among the work classes except for self-emp-inc and federal government. Federal government is seen as the most elite in the public sector, which most likely explains the higher chance of earning more than 50,000.


1.4.2 occupation

```
[19]: sns.histplot(x=df['occupation'], hue=df['salary'], color='green', bins=15)
plt.tight_layout()
plt.xticks(rotation = 50)
plt.show()
```




1.4.3 race

```
[26]: sns.histplot(x = df['race'], hue = df['salary'], color = 'green', bins = 15,
    ↪, palette = ['#a4def5', 'red'])
plt.tight_layout()
plt.xticks(rotation = 50)
plt.show()
```



1.4.4 sex

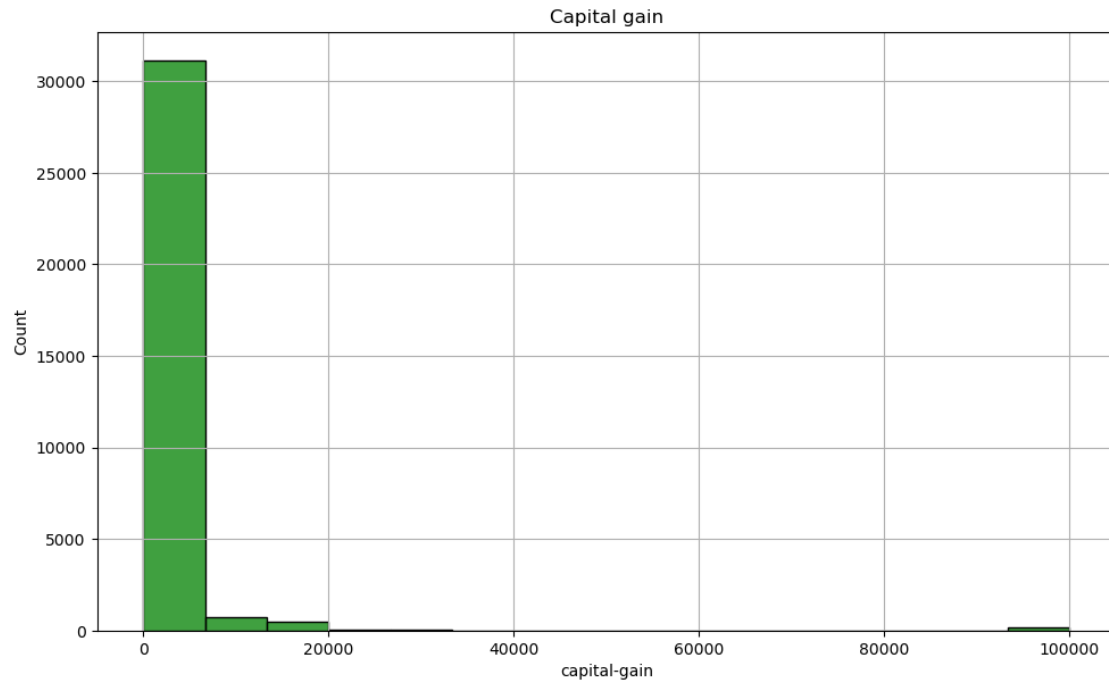
```
[27]: plt.figure(figsize = [8,4])
sns.countplot(y=df['sex'],hue = df['salary'],palette=['#a4def5','#e1a4f5'])
plt.tight_layout()
plt.show()
```



- The percentage of males who make greater than 50,000 is much greater than the percentage of females that make the same amount. This will certainly be a significant factor, and should be a feature considered in our prediction model.


1.4.5 capital gain

```
[28]: plt.figure(figsize=(10,6))
sns.histplot(df['capital-gain'], color = 'green', bins = 15)
plt.tight_layout()
plt.grid(True)
plt.title('Capital gain')
plt.show()
```



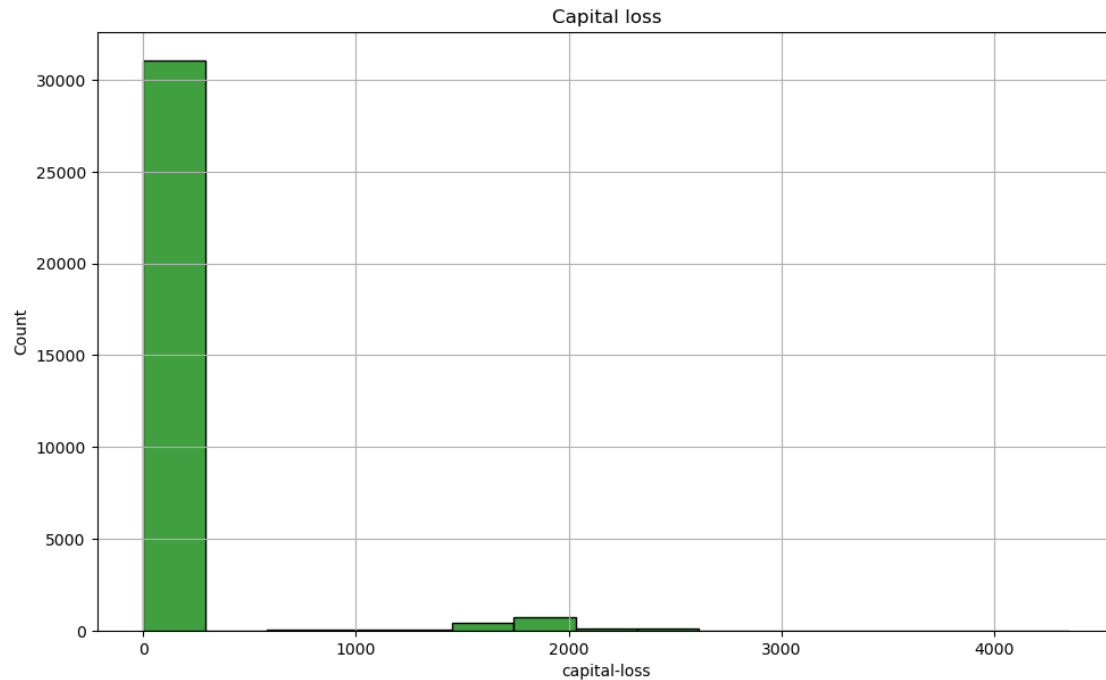
- From the graph we can see that the distribution of capital gain is very skewed.
- And there are outliers at data poitn 100000.

```
[29]: sns.histplot(x=df['capital-gain'],hue=df['salary'],color='green',bins=15)  
plt.tight_layout()  
plt.show()
```




1.4.6 capital loss

```
[30]: plt.figure(figsize=(10,6))
sns.histplot(df['capital-loss'],color='green',bins=15)
plt.tight_layout()
plt.grid(True)
plt.title('Capital loss')
plt.show()
```



1.4.7 hours per week

```
[31]: sns.histplot(df['hours-per-week'],color='green',bins=15)
plt.tight_layout()
plt.grid(True)
plt.show()
```




- We can see that vast majority of values are of 40 hours. what we can do is make 3 classes, i.e <40hrs, 40hrs and >40hrs, and check whether it is significant or not.

```
[32]: def hrs_edit(val):
      if (val<40):
          return ('<40 hrs')
      elif (val==40):
          return ('40 hrs')
      else:
          return ('>40hrs')
```

```
[33]: df['hours-per-week']=df['hours-per-week'].apply(hrs_edit)
```

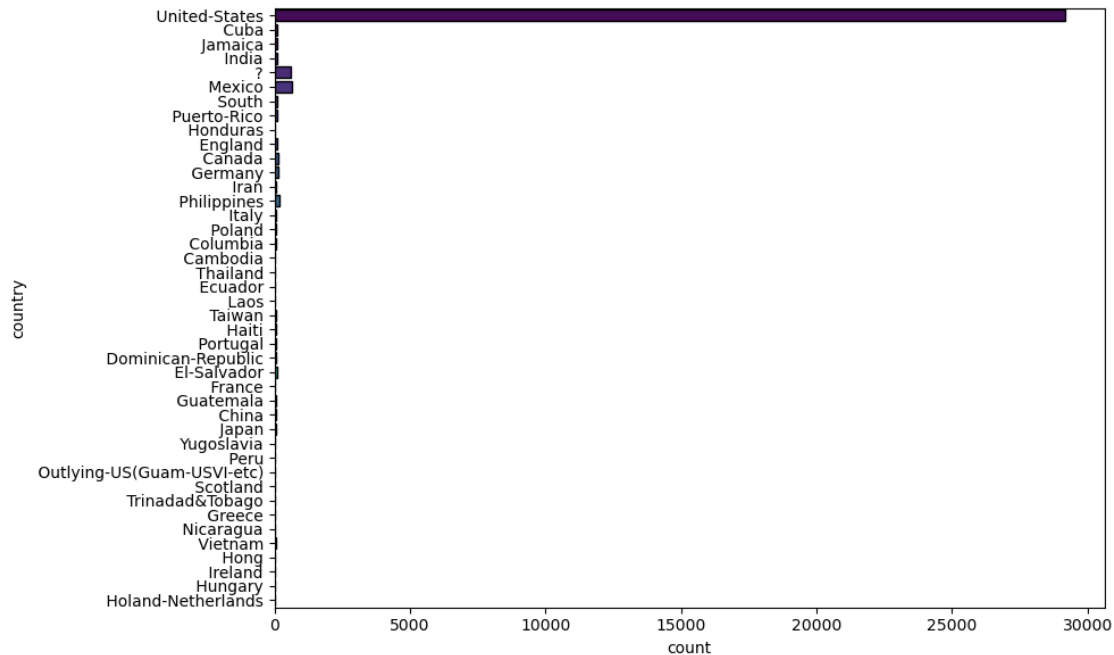
```
[34]: sns.countplot(x=df['hours-per-week'],hue=df['salary'],palette='viridis',
                    saturation=0.9,edgecolor="black",order=['<40 hrs','40_
                    ↪hrs','>40hrs'])
plt.tight_layout()
plt.show()
```

- The percentage of individuals making over 50,000 drastically decreases when less than 40 hours per week, and increases significantly when greater than 40 hours per week.

1.4.8 country

```
[35]: plt.figure(figsize=(10,6))
sns.countplot(y=df['country'],palette='viridis',saturation=0.
↪9,edgecolor="black")
plt.tight_layout()
plt.show()
```



1.4.9 Feature engineering

1. We will drop the features:
 - fnlwgt
 - education
 - relationship
 - race
2. Impute Nan values with mode.
3. train - test split.
4. Lable encoding.

1.4.10 Dropping Education - Education Num is enough.

1.4.11 Dropping Final Weight - highly discrete data so not useful

```
[36]: df = df.drop(['education', 'fnlwgt', 'race', 'relationship'], axis = 1)
df
```

```
[36]:
```

	age	workclass	education-num	marital-status	\
0	39	State-gov	13	Never-married	
1	50	Self-emp-not-inc	13	Married-civ-spouse	
2	38	Private	9	Divorced	
3	53	Private	7	Married-civ-spouse	
4	28	Private	13	Married-civ-spouse	

```

...      ...      ...      ...      ...
32556    27      Private      12    Married-civ-spouse
32557    40      Private      9    Married-civ-spouse
32558    58      Private      9      Widowed
32559    22      Private      9    Never-married
32560    52    Self-emp-inc      9    Married-civ-spouse

      occupation      sex      capital-gain      capital-loss      hours-per-week \
0      Adm-clerical      Male      2174      0      40 hrs
1      Exec-managerial      Male      0      0      <40 hrs
2      Handlers-cleaners      Male      0      0      40 hrs
3      Handlers-cleaners      Male      0      0      40 hrs
4      Prof-specialty      Female      0      0      40 hrs
...      ...      ...      ...      ...
32556      Tech-support      Female      0      0      <40 hrs
32557    Machine-op-inspct      Male      0      0      40 hrs
32558      Adm-clerical      Female      0      0      40 hrs
32559      Adm-clerical      Male      0      0      <40 hrs
32560      Exec-managerial      Female      15024      0      40 hrs

      country      salary
0      United-States      <=50K
1      United-States      <=50K
2      United-States      <=50K
3      United-States      <=50K
4      Cuba      <=50K
...      ...      ...
32556    United-States      <=50K
32557    United-States      >50K
32558    United-States      <=50K
32559    United-States      <=50K
32560    United-States      >50K

```

[32561 rows x 11 columns]

1.4.12 Replacing ? with NaN

```
[37]: df.isin(['?']).sum()
```

```

[37]: age      0
      workclass      1835
      education-num      0
      marital-status      0
      occupation      1843
      sex      0
      capital-gain      0
      capital-loss      0

```

```
hours-per-week      0
country             583
salary              0
dtype: int64
```

```
[38]: df['workclass'].replace(' ?',0,inplace = True)
df['occupation'].replace(' ?',0,inplace = True)
df['country'].replace(' ?',0,inplace = True)
```

```
[39]: df['workclass'].replace(0,np.nan,inplace = True)
df['occupation'].replace(0,np.nan,inplace = True)
df['country'].replace(0,np.nan,inplace = True)
```

```
[40]: df["workclass"] = df["workclass"].fillna(df["workclass"].mode()[0])
df["occupation"] = df["occupation"].fillna(df["occupation"].mode()[0])
df["country"] = df["country"].fillna(df["country"].mode()[0])
```

```
[41]: df["workclass"].value_counts()
```

```
[41]: Private          24532
Self-emp-not-inc    2541
Local-gov           2093
State-gov           1298
Self-emp-inc        1116
Federal-gov         960
Without-pay         14
Never-worked         7
Name: workclass, dtype: int64
```

```
[42]: df['marital-status'].unique()
```

```
[42]: array([' Never-married', ' Married-civ-spouse', ' Divorced',
        ' Married-spouse-absent', ' Separated', ' Married-AF-spouse',
        ' Widowed'], dtype=object)
```

```
[43]: def married(val):
    if val==' Never-married':
        return 'not-married'
    elif val==' Divorced':
        return 'not-married'
    elif val==' Separated':
        return 'not-married'
    elif val==' Widowed':
        return 'not-married'
    else:
        return 'married'
```

```
[44]: df['marital-status']=df['marital-status'].apply(married)
```

```
[45]: df['marital-status'].unique()
```

```
[45]: array(['not-married', 'married'], dtype=object)
```

1.5 Label Encoding

```
[46]: encoder = preprocessing.LabelEncoder()
```

```
[47]: df['workclass'] = encoder.fit_transform(df['workclass'])
df['marital-status'] = encoder.fit_transform(df['marital-status'])
df['occupation'] = encoder.fit_transform(df['occupation'])
df['sex'] = encoder.fit_transform(df['sex'])
df['country'] = encoder.fit_transform(df['country'])
df['salary'] = encoder.fit_transform(df['salary'])
df['hours-per-week'] = encoder.fit_transform(df['hours-per-week'])
```

```
[48]: df
```

```
[48]:
```


	age	workclass	education-num	marital-status	occupation	sex	\
0	39	6	13	1	0	1	
1	50	5	13	0	3	1	
2	38	3	9	1	5	1	
3	53	3	7	0	5	1	
4	28	3	13	0	9	0	
...	
32556	27	3	12	0	12	0	
32557	40	3	9	0	6	1	
32558	58	3	9	1	0	0	
32559	22	3	9	1	0	1	
32560	52	4	9	0	3	0	

	capital-gain	capital-loss	hours-per-week	country	salary
0	2174	0	0	38	0
1	0	0	1	38	0
2	0	0	0	38	0
3	0	0	0	38	0
4	0	0	0	4	0
...
32556	0	0	1	38	0
32557	0	0	0	38	1
32558	0	0	0	38	0
32559	0	0	1	38	0
32560	15024	0	0	38	1

```
[32561 rows x 11 columns]
```

```
[49]: sns.pairplot(df,hue=None)
```

```
[49]: <seaborn.axisgrid.PairGrid at 0x2054bbfc580>
```



1.5.1 Model train and test

1.6 Logistic Regression

```
[50]: logistic = LogisticRegression()
```

```
[51]: X = df.iloc[:, :-1]  
y = df.iloc[:, -1]
```

```
[52]: X_train,X_test,y_train,y_test = model_selection.train_test_split(X,y)
      logistic.fit(X_train,y_train)
```

```
C:\ProgramData\Anaconda3\lib\site-
packages\sklearn\linear_model\_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
[52]: LogisticRegression()
```

```
[65]: y_pred = logistic.predict(X_test)
```

```
[66]: print(confusion_matrix(y_test,y_pred))
```

```
[[5726  430]
 [1108  877]]
```

```
[67]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.84	0.93	0.88	6156
1	0.67	0.44	0.53	1985
accuracy			0.81	8141
macro avg	0.75	0.69	0.71	8141
weighted avg	0.80	0.81	0.80	8141

```
[68]: print("Accuracy",round(accuracy_score(y_test, y_pred)*100),"%")
```

Accuracy 81 %

1.7 Random Forest

```
[57]: random_forest = RandomForestClassifier(n_estimators=10,
      random_state=0)
      random_forest.fit(X_train, y_train)
      Y_prediction = random_forest.predict(X_test)
      random_forest.score(X_train, y_train)
      acc_random_forest = round(random_forest.score(X_train, y_train) * 100, 2)
```

```
[58]: print(confusion_matrix(y_test,Y_prediction))
```

```
[[5627  529]
 [ 779 1206]]
```

```
[59]: print(classification_report(y_test,Y_prediction))
```

	precision	recall	f1-score	support
0	0.88	0.91	0.90	6156
1	0.70	0.61	0.65	1985
accuracy			0.84	8141
macro avg	0.79	0.76	0.77	8141
weighted avg	0.83	0.84	0.84	8141

```
[60]: print("Accuracy",round(accuracy_score(y_test, Y_prediction)*100),"%")
```

Accuracy 84 %

1.8 Decision tree

```
[61]: classifier= DecisionTreeClassifier(criterion='entropy', random_state=10)
classifier.fit(X_train, y_train)
```

```
[61]: DecisionTreeClassifier(criterion='entropy', random_state=10)
```

```
[62]: y_pred = classifier.predict(X_test)
```

```
[64]: print("Accuracy", round(accuracy_score(y_test, y_pred)*100), "%")
```

Accuracy 82 %

```
[ ]:
```